

Julia

Not a simple language.

Strengths

- performance
- parallelism out of the box
- scientific computing (*a lot* of amenities)

Typing

- “dynamically”-typed
- optionally typed with type inference

Overview

- `x = 10`
- meta programming (macros)
- functional programming: `id = a -> a`
- bang! convention
- usual PL features (loops, functions, etc.)

The Amenities

- unicode
- `Int8`, `UInt8`, ..., `Int128`, `UInt128`, `Float16` (32 & 64)
- `1 + 2im` (complex numbers)
- `-4 // -12 == 1 // 3` (rational numbers)
- binary (`0b10`), octal (`0o010`), hexa (`0xA`) values
- chaining comparisons: `0 <= a < 1`
- *much* more...

```

1  #  $\Sigma(\backslash\text{Sigma})$ 
2
3  # nice and readable
4  function  $\Sigma 1(f, n::\text{Int}, m::\text{Int})::\text{Int}$ 
5      sum = 0
6      for i = n:m
7          sum += f(i)
8      end
9      return sum
10 end
11
12 n = 1
13 m = 3
14 f = a -> a^2
15 result =  $\Sigma 1(f, n, m)$ 
16
17 print("Normal ( $\Sigma 1$ ): ")
18 println(result)
19
20
21 # weird and compact
22  $\Sigma 2(f, n, m) = n \leq m ? f(n) + \Sigma 2(f, n + 1, m) : 0$ 
23
24 arguments = (f, n, m)
25 result =  $\Sigma 2(\text{arguments}...)$ 
26
27 print("Recursive ( $\Sigma 2$ ): ")
28 println(result)
29

```

Parallelism

- Julia Coroutines (Green Threading)
- Multi-Threading (Experimental)
- Multi-Core or Distributed Processing

Coroutines

```
1  # tasks (symmetric coroutines)
2  println("-- Tasks")
3
4  producer = @task begin
5      for i = 1:3
6          yieldto(consumer, i * 2)
7      end
8  end
9
10 consumer = @task begin
11     while true
12         println(yieldto(producer))
13     end
14 end
15
16 schedule(producer)
17 schedule(consumer)
18 wait(producer)
19
```

```
1  # channels
2  println("-- Channels")
3
4  function producer(c)
5      for i = 1:5
6          put!(c, i * 2)
7      end
8      close(c)
9  end
10
11  c = Channel{Int}(1)
12  @async producer(c)
13
14  for i in c
15      println(i)
16  end
17
```

Other Task Functions

- **yield()** : yields to the scheduler. A task that calls this function is still runnable, and will be restarted immediately if there are no other runnable tasks.
- **wait(task)** : blocks until the given task is done. (wait behaves differently depending on its argument's type.)
- **fetch(channel)** : waits until there is an item in the channel and returns the item without removing it.
- **notify(condition, ...)** : wake up tasks waiting for a condition.
- **bind(channel, task)** : associates the lifetime of channel with a task.

Multi-Threading

- likely to change
- starts up with `JULIA_NUM_THREADS` (default 1)
- synchronization primitives (locks, semaphores, etc) are available
- functions `Threads.nthreads` and `Threads.threadid`
- ```
@threads for i = 1:10
 array[i] = Threads.threadid()
end
```
- atomicity:  

```
a = Threads.Atomic{Int}(15)
Threads.atomic_add!(a, 5)
```

# Distributed Processing

- `julia -p n` (starts with  $n$  local worker processes) or use the `--machine-file` option to start Julia processes on other machines
- `Distributed` module within the standard library
- message passing through remote calls
- remote references (`Future` and `RemoteChannel`)

# Distributed Processing

- remote calls return a Future (immediately)  

```
sum = remotecall(+, myid(), 2, 3)
fetch(sum) # 5
same as remotecall_fetch(+, myid(), 2, 3)
```
- remotecall is low-level, best to use @spawn most of the time  

```
sum = @spawn 2 + 3
fetch(sum)
```
- @async is similar to @spawn, but only runs tasks on the local process
- a normal Channel can't be shared between workers, hence the RemoteChannel

# Distributed Processing

- `@distributed (reduce)`  
    `@distributed (+) for i = 1:10 i end`  
    `# result is 55`
- without the reduce function, `@distributed` executes asynchronously and returns an array of futures immediately without waiting completion
- `@everywhere` (all processes)  
    `@everywhere include("SuperDuperModule.jl")`  
    `# loads the module on all processes`
- `@sync` synchronises enclosed `@async`, `@spawn`, `@spawnat`  
    `@distributed` macros
- `pmap` function, just like `map`, but parallel



# Distributed Processing

- MPI

```
import MPI
```

- a Julia wrapper of the MPI protocol
- `mpirun -np 4 ./julia example.jl`

**parallel depth-first scheduling**

?